


Postmodern for GPUs



Conor Hoekstra



code_report



codereport



Programming Language Rankings (2025 Aug)



by code_report





☒ StackOverflow ☒ Octoverse ☒ RedMonk ☒ Languish ☐ JetBrains
☐ IEEE Spectrum ☐ PYPL ☐ TIOBE ☐ GitHub 2.0 [Rankings Overview](#)

☒ Exclude "Edge Languages" | Number of Languages:
Months for Delta (Δ): | Languages

	Language	Avg	StDev	n ¹	3m Δ		Language	Avg	StDev	n ¹	3m Δ
1	JavaScript	1.5	0.57	4	-	11	Rust	12.66	3.21	3	-
2	Python	2	1.41	4	-	12	Kotlin	14	1.73	3	-
3	TypeScript	4.5	1.73	4	-	13	Swift	15.66	5.13	3	-
4	Java	4.5	1.73	4	-	14	PowerShell	15.66	6.65	3	(1)
5	C#	5.75	1.5	4	-	15	R	16	5	3	(1)
6	C++	7	1.41	4	-	16	Dart	17.66	2.3	3	-
7	PHP	8.5	3.69	4	-	17	Ruby	18	10	3	-
8	Shell	9.5	3.69	4	(1)	18	Lua	21.5	7.77	2	-
9	C	9.75	0.95	4	(1)	19	Objective-C	23	14.14	2	(2)
10	Go	10.25	2.5	4	-	20	VBA	25	1.41	2	(1)

1 - The number of (selected) ranking websites this language shows up in.

If you have suggestions or find a bug, you can open an [issue](#) here.



About Me

Conor Hoekstra / @code_report





383 Videos



48 (32) Talks

Algorithms +
Data
Structures =
Programs

252 Episodes
@adspthepodcast



Array Cast

115 Episodes
@arraycast



TACIT
≡ ∘ ^ ∨ ∘ ≡
TALK

27 Episodes
@codereport





Premium

Search



Home



Shorts



Subscriptions



YouTube Mu...



You



Downloads



code_report

@code_report · 60.8K subscribers · 352 videos

Welcome to the code_report YouTube channel. ...more

twitter.com/code_report and 3 more links

Customize channel

Manage videos

Home

Videos

Live

Podcasts

Playlists

Community



Latest


Popular

Oldest



1 Problem, 24 Programming Languages

375K views · 1 year ago



From C → C++ → Rust

170K views · 1 year ago




1 Problem, 16 Programming Languages
(C++ vs Rust vs Haskell vs Python vs APL...)

158K views · 3 years ago




Functional vs Array Programming

131K views · 3 years ago



<https://github.com/codereport/Content>

C++ CORE COMPUTE LIBRARIES



The Standard Library for your entire system

<https://github.com/NVIDIA/libcudacxx>



The C++ parallel algorithms library

<https://github.com/NVIDIA/thrust>



Cooperative primitives for CUDA C++

<https://github.com/NVIDIA/cub>

CUDA CORE COMPUTE LIBRARIES



The Standard Library for your entire system

<https://github.com/NVIDIA/tremailibcudacxx>



The C++ parallel algorithms library

<https://github.com/NVIDIA/thrust>



Cooperative primitives for CUDA C++

<https://NVIDIA.cccl/tremain/cub>



Overview

Problem 1




API Overview



Problem 2



Sum of Squares (SOS)





```
#include <iostream>
#include <numeric>
#include <ranges>

auto sos(int N) {
    return std::ranges::views::iota(0, N) |
           std::ranges::views::transform([](int x) { return x * x; }) |
           std::ranges::fold_left(0, std::plus{});
}

int main() { std::cout << sos(10) << std::endl; }
```



```
#include <algorithm>
#include <iostream>
#include <numeric>
#include <ranges>

auto sos(int N) {
    return std::ranges::views::iota(0, N) |
           std::ranges::views::transform([](int x) { return x * x; }) |
           std::ranges::fold_left(0, std::plus{});
}

int main() { std::cout << sos(10) << std::endl; }
```



```
#include <algorithm>
#include <iostream>
#include <ranges>

auto sos(int N) {
    return std::ranges::views::iota(0, N) |
           std::ranges::views::transform([](int x) { return x * x; }) |
           std::ranges::fold_left(0, std::plus{});
}

int main() { std::cout << sos(10) << std::endl; }
```




```
#include <algorithm>
#include <iostream>
#include <ranges>

auto sos(int N) {
    return std::ranges::fold_left(
        std::ranges::views::iota(0, N) |
        std::ranges::views::transform([](int x) { return x * x; }),
        0,
        std::plus{});
}

int main() { std::cout << sos(10) << std::endl; }
```



```
#include <algorithm>
#include <print>
#include <ranges>

auto sos(int N) {
    return std::ranges::fold_left(
        std::ranges::views::iota(0, N) |
        std::ranges::views::transform([](int x) { return x * x; }),
        0,
        std::plus{});
}

int main() { std::print("{} ", sos(10)); }
```





```
#include <thrust/device_vector.h>
#include <thrust/functional.h>
#include <thrust/host_vector.h>
#include <thrust/reduce.h>
#include <thrust/sequence.h>
#include <thrust/transform.h>
#include <iostream>

// Functor to square a number
struct square {
    __host__ __device__ int operator()(const int& x) const { return x * x; }
};

auto sos(int N) {
    // Create a device vector and fill it with sequence 0, 1, 2, ..., N-1
    thrust::device_vector<int> d_vec(N);
    thrust::sequence(d_vec.begin(), d_vec.end());

    // Square each element and sum the result
    return thrust::transform_reduce(
        d_vec.begin(), d_vec.end(), square(), 0, thrust::plus<int>());
}

int main() {
    auto result = sos(10);
    std::cout << result << std::endl;
    return 0;
}
```



```
#include <thrust/device_vector.h>
#include <thrust/functional.h>
#include <thrust/host_vector.h>
#include <thrust/reduce.h>
#include <thrust/sequence.h>
#include <thrust/transform_reduce.h>
#include <iostream>

// Functor to square a number
struct square {
    __host__ __device__ int operator()(const int& x) const { return x * x; }
};

auto sos(int N) {
    // Create a device vector and fill it with sequence 0, 1, 2, ..., N-1
    thrust::device_vector<int> d_vec(N);
    thrust::sequence(d_vec.begin(), d_vec.end());

    // Square each element and sum the result
    return thrust::transform_reduce(
        d_vec.begin(), d_vec.end(), square(), 0, thrust::plus<int>());
}

int main() {
    auto result = sos(10);
    std::cout << result << std::endl;
    return 0;
}
```



```
#include <thrust/device_vector.h>
#include <thrust/sequence.h>
#include <thrust/transform_reduce.h>
#include <iostream>

// Functor to square a number
struct square {
    __host__ __device__ int operator()(const int& x) const { return x * x; }
};

auto sos(int N) {
    // Create a device vector and fill it with sequence 0, 1, 2, ..., N-1
    thrust::device_vector<int> d_vec(N);
    thrust::sequence(d_vec.begin(), d_vec.end());

    // Square each element and sum the result
    return thrust::transform_reduce(
        d_vec.begin(), d_vec.end(), square(), 0, thrust::plus<int>());
}

int main() {
    auto result = sos(10);
    std::cout << result << std::endl;
    return 0;
}
```




```
#include <thrust/device_vector.h>
#include <thrust/sequence.h>
#include <thrust/transform_reduce.h>
#include <iostream>

struct square {
    __host__ __device__ int operator()(const int& x) const { return x * x; }
};

auto sos(int N) {
    thrust::device_vector<int> d_vec(N);
    thrust::sequence(d_vec.begin(), d_vec.end());
    return thrust::transform_reduce(
        d_vec.begin(), d_vec.end(), square(), 0, thrust::plus<int>());
}

int main() {
    auto result = sos(10);
    std::cout << result << std::endl;
    return 0;
}
```



```
#include <thrust/device_vector.h>
#include <thrust/sequence.h>
#include <thrust/transform_reduce.h>
#include <iostream>

auto sos(int N) -> int {
    thrust::device_vector<int> d_vec(N);
    thrust::sequence(d_vec.begin(), d_vec.end());
    return thrust::transform_reduce(
        d_vec.begin(),
        d_vec.end(),
        [] __host__ __device__(int x) { return x * x; },
        0,
        thrust::plus<int>());
}

int main() {
    auto result = sos(10);
    std::cout << result << std::endl;
    return 0;
}
```



```
#include <thrust/device_vector.h>
#include <thrust/sequence.h>
#include <thrust/transform_reduce.h>
#include <iostream>

auto sos(int N) -> int {
    auto iota = thrust::counting_iterator<int>(0);
    return thrust::transform_reduce(
        iota,
        iota + N,
        [] __host__ __device__(int x) { return x * x; },
        0,
        thrust::plus<int>());
}

int main() {
    auto result = sos(10);
    std::cout << result << std::endl;
    return 0;
}
```



```
#include <thrust/iterator/counting_iterator.h>
#include <thrust/iterator/transform_iterator.h>
#include <thrust/reduce.h>
#include <iostream>

auto sos(int N) -> int {
    auto iota = thrust::counting_iterator<int>(0);
    auto map = thrust::make_transform_iterator(
        iota, [] __host__ __device__(int x) { return x * x; });
    return thrust::reduce(map, map + N, 0);
}

int main() {
    auto result = sos(10);
    std::cout << result << std::endl;
    return 0;
}
```





Parrot



Parrot

A high level, parallel, array-based library
with implicit fusion



```
auto sos(int N) {  
  
}
```



```
auto sos(int N) {  
    return parrot::range(N)  
}
```



```
auto sos(int N) {  
    return parrot::range(N)  
        .sq()  
}
```



```
auto sos(int N) {  
    return parrot::range(N)  
        .sq()  
        .sum();  
}
```



```
#include "parrot.hpp"
```

```
auto sos(int N) {  
    return parrot::range(N)  
        .sq()  
        .sum();  
}
```

```
int main() { sos(10).print(); }
```




```
#include "parrot.hpp"
```

```
auto sos(int N) {  
    return parrot::range(N).sq().sum();  
}
```

```
int main() { sos(10).print(); }
```

```
nvcc  
--extended-lambda  
-std=c++20  
-arch=native  
-I ~/parrot-cpp  
file_name.cu
```



Parrot API

Fused
1-index Maps (Unary)

★ [map](#)
[abs](#)
[dble](#)
[enumerate](#)
[even](#)
[half](#)
[log](#)
[exp](#)
[neg](#)
[odd](#)
[rand](#)
[sign](#)
[sq](#)
[sqrt](#)

1-index Maps (Binary)

★ [map2](#)
[add \(+\)](#)
[div \(/\)](#)
[gt \(>\)](#)
[gte \(>=\)](#)
[idiv](#)
[lt \(<\)](#)
[lte \(<=\)](#)
[max](#)
[min](#)
[minus \(-\)](#)
[times \(*\)](#)
[eq \(==\)](#)
[pairs](#)

2-index Maps

★ [map_adj](#)
[deltas](#)
[differ](#)
Joins
[append](#)
[prepend](#)
Products
[cross](#)
[outer](#)
Reshapes
[take](#)
[drop](#)
[transpose](#)
[reshape](#)
[cycle](#)
[repeat](#)
Copying
[replicate](#)
Permutations
[rev](#)
[gather](#)

Conditionally Fused

Compactions
★ [keep](#)
[filter](#)
[where](#)
[uniq](#)
[distinct](#)

Materializing
Reductions

★ [reduce](#)
[all](#)
[any](#)
[maxr](#)
[max_by_key](#)
[minr](#)
[minmax](#)
[prod](#)
[sum](#)
Scans
★ [scan](#)
[alls](#)
[anys](#)
[maxs](#)
[mins](#)
[prods](#)
[sums](#)

Permutations

[sort](#)
[sort_by](#)
[sort_by_key](#)

Compactions

[rle](#)

Copying

[replicate](#)

Split-Reductions

[chunk_by_reduce](#)

Comparisons

[match](#)

Properties

[size](#)
[rank](#)
[shape](#)

Accessors

[value](#)
[front](#)
[back](#)
[to_host](#)

Array Creation

[array](#)
[range](#)
[scalar](#)
[matrix](#)

I/O

[print](#)

Function Objects

Accessors

[fst](#)
[snd](#)

Binary Operations

[eq](#)
[gt](#)
[gte](#)
[lt](#)
[lte](#)
[max](#)
[min](#)
[mul](#)
[add](#)

Problem 2

(taken from PaddlePaddle)



```
// https://github.com/PaddlePaddle/Paddle/blob/80f1123eb0c...

template <typename T>
static void Get_____(
    const phi::GPUContext& dev_ctx,
    const DenseTensor* input_tensor,
    const int64_t num_cols,
    const int64_t num_rows,
    T* out_tensor,
    int64_t* indices_tensor)
{
    DenseTensor input_tmp;
    input_tmp.Resize(common::make_ddim({num_rows, num_cols}));
    T* input_tmp_data = dev_ctx.Alloc<T>(&input_tmp);
    phi::Copy(dev_ctx, *input_tensor, dev_ctx.GetPlace(), false, &input_tmp);

    thrust::device_ptr<T> out_tensor_ptr(out_tensor);
    thrust::device_ptr<int64_t> indices_tensor_ptr(indices_tensor);

    for (int64_t i = 0; i < num_rows; ++i) {
        T* begin = input_tmp_data + num_cols * i;
        T* end = input_tmp_data + num_cols * (i + 1);
        thrust::device_vector<int64_t> indices_data(num_cols);
        thrust::sequence(thrust::device,
            indices_data.begin(),
            indices_data.begin() + num_cols);
        thrust::sort_by_key(thrust::device, begin, end, indices_data.begin());
        int unique = 1 + thrust::inner_product(thrust::device,
            begin,
            end - 1,
            begin + 1,
            0,
            thrust::plus<int>(),
            thrust::not_equal_to<T>());
        thrust::device_vector<T> keys_data(unique);
        thrust::device_vector<int64_t> cnts_data(unique);
        thrust::reduce_by_key(thrust::device,
            begin,
            end,
            thrust::constant_iterator<int>(1),
            keys_data.begin(),
            cnts_data.begin());
        auto it = thrust::max_element(
            thrust::device, cnts_data.begin(), cnts_data.begin() + unique);
        T ____ = keys_data[it - cnts_data.begin()];
        int64_t counts = cnts_data[it - cnts_data.begin()];
        auto pos = thrust::find(thrust::device, begin, end, mode);
        int64_t index = indices_data[pos - begin + counts - 1];
        out_tensor_ptr[i] = static_cast<T>(mode);
        indices_tensor_ptr[i] = static_cast<int64_t>(index);
    }
}
```



```
// https://github.com/PaddlePaddle/Paddle/blob/80f1123eb0c...
```

```
template <typename T>
static void Get_____(/* ... */) {
    // initialization

    for (int64_t i = 0; i < num_rows; ++i) {
        T* begin = input_tmp_data + num_cols * i;
        T* end = input_tmp_data + num_cols * (i + 1);
        thrust::device_vector<int64_t> indices_data(num_cols);
        thrust::sequence(thrust::device,
            indices_data.begin(),
            indices_data.begin() + num_cols);
        thrust::sort_by_key(thrust::device, begin, end, indices_data.begin());
        int unique = 1 + thrust::inner_product(thrust::device,
            begin,
            end - 1,
            begin + 1,
            0,
            thrust::plus<int>(),
            thrust::not_equal_to<T>());
        thrust::device_vector<T> keys_data(unique);
        thrust::device_vector<int64_t> cnts_data(unique);
        thrust::reduce_by_key(thrust::device,
            begin,
            end,
            thrust::constant_iterator<int>(1),
            keys_data.begin(),
            cnts_data.begin());
        auto it = thrust::max_element(
            thrust::device, cnts_data.begin(), cnts_data.begin() + unique);
        T ____ = keys_data[it - cnts_data.begin()];
        int64_t counts = cnts_data[it - cnts_data.begin()];
        auto pos = thrust::find(thrust::device, begin, end, mode);
        int64_t index = indices_data[pos - begin + counts - 1];
        out_tensor_ptr[i] = static_cast<T>(mode);
        indices_tensor_ptr[i] = static_cast<int64_t>(index);
    }
}
```



```
int unique = 1 + thrust::inner_product(
    begin,
    end - 1,
    begin + 1,
    0,
    thrust::plus<int>(),
    thrust::not_equal_to<T>());
```




```
int unique = 1 + thrust::inner_product(  
    begin,  
    end - 1,  
    begin + 1,  
    0,  
    thrust::plus<int>(),  
    thrust::not_equal_to<T>());
```



```
int unique = thrust::unique_count(begin, end);
```

[[digression]]

[illegible]



```
template <typename I>
auto unique_count(I first, I last) {
    auto zip = thrust::make_zip_iterator(first, first + 1);
    auto map = thrust::make_transform_iterator(
        zip, thrust::make_zip_function(thrust::not_equal_to{}));
    return 1 + thrust::reduce(map,
                              map + (last - first - 1),
                              0,
                              thrust::plus{});
}
```



```
template <class Derived,
          class ForwardIt,
          class BinaryPred>
typename thrust::iterator_traits<ForwardIt>::difference_type
_CCCL_HOST_DEVICE
unique_count(execution_policy<Derived> &policy,
             ForwardIt first,
             ForwardIt last,
             BinaryPred binary_pred)
{
    if (first == last) {
        return 0;
    }
    auto size = thrust::distance(first, last);
    auto it    = thrust::make_zip_iterator(thrust::make_tuple(first, thrust::next(first)));
    return 1 + thrust::count_if(policy, it, thrust::next(it, size - 1),
                               zip_adj_not_predicate<BinaryPred>{binary_pred});
}
```



```
template <typename I>
auto unique_count_zip_count_if(I first, I last) {
    auto zip = thrust::make_zip_iterator(first, first + 1);
    auto neq = thrust::make_zip_function(thrust::not_equal_to{});
    return 1 + thrust::count_if(zip, zip + (last - first - 1), neq);
}
```



```
auto unique_count(auto data) {  
    return data.map_adj(parrot::neq{}).sum() + 1;  
}
```




```
auto unique_count(auto data) {  
    return data.differ().sum() + 1;  
}
```



`thrust::transform_reduce`

`transform_iterator + reduce`

★ `map + reduce`

`thrust::unique_count`

`zip_iterator + count_if`

★ `map_adj + sum`

`thrust::tabulate`

`counting_iterator +
transform`

`range + ★ map`

[[end of digression]]



```
// https://github.com/PaddlePaddle/Paddle/blob/80f1123eb0c...
```

```
template <typename T>
static void Get_____(/* ... */) {
    // initialization

    for (int64_t i = 0; i < num_rows; ++i) {
        T* begin = input_tmp_data + num_cols * i;
        T* end = input_tmp_data + num_cols * (i + 1);
        thrust::device_vector<int64_t> indices_data(num_cols);
        thrust::sequence(thrust::device,
            indices_data.begin(),
            indices_data.begin() + num_cols);
        thrust::sort_by_key(thrust::device, begin, end, indices_data.begin());
        int unique = thrust::unique_count(thrust::device, begin, end);
        thrust::device_vector<T> keys_data(unique);
        thrust::device_vector<int64_t> cnts_data(unique);
        thrust::reduce_by_key(thrust::device,
            begin,
            end,
            thrust::constant_iterator<int>(1),
            keys_data.begin(),
            cnts_data.begin());
        auto it = thrust::max_element(
            thrust::device, cnts_data.begin(), cnts_data.begin() + unique);
        T ____ = keys_data[it - cnts_data.begin()];
        int64_t counts = cnts_data[it - cnts_data.begin()];
        auto pos = thrust::find(thrust::device, begin, end, mode);
        int64_t index = indices_data[pos - begin + counts - 1];
        out_tensor_ptr[i] = static_cast<T>(mode);
        indices_tensor_ptr[i] = static_cast<int64_t>(index);
    }
}
```



```
// https://github.com/PaddlePaddle/Paddle/blob/80f1123eb0c...
```

```
template <typename T>
static void Get_____(/* ... */) {
    // initialization

    for (int64_t i = 0; i < num_rows; ++i) {
        T* begin          = input_tmp_data + num_cols * i;
        T* end            = input_tmp_data + num_cols * (i + 1);
        auto indices_data = thrust::device_vector<int64_t>(num_cols);

        thrust::sequence(indices_data.begin(), indices_data.begin() + num_cols);
        thrust::sort_by_key(begin, end, indices_data.begin());

        int unique      = thrust::unique_count(thrust::device, begin, end);
        auto keys_data   = thrust::device_vector<T>(unique);
        auto cnts_data   = thrust::device_vector<int64_t>(unique);

        thrust::reduce_by_key(
            begin,
            end,
            thrust::constant_iterator<int>(1),
            keys_data.begin(),
            cnts_data.begin());

        auto it = thrust::max_element(cnts_data.begin(), cnts_data.begin() + unique);

        T ____          = keys_data[it - cnts_data.begin()];
        int64_t counts    = cnts_data[it - cnts_data.begin()];
        auto pos          = thrust::find(begin, end, mode);
        int64_t index     = indices_data[pos - begin + counts - 1];
        out_tensor_ptr[i] = static_cast<T>(mode);
        indices_tensor_ptr[i] = static_cast<int64_t>(index);
    }
}
```



```
template <typename Array>
auto GetModeBySort_Parrot(
    const Array& data, int num_rows, int num_cols) {
    using T = typename Array::value_type;
    std::vector<thrust::pair<T, int>> results;

    for (int r = 0; r < num_rows; ++r) {
        auto mode = parrot::stats::mode(data.row(r)).value();
        auto index = data.row(r).last_index_of(mode);
        results.push_back(thrust::make_pair(mode, index));
    }

    return parrot::array(results);
}
```



```
template <typename Array>
auto GetModeBySort_Parrot(
    const Array& data, int num_rows, int num_cols) {
    using T = typename Array::value_type;
    std::vector<thrust::pair<T, int>> results;

    for (int r = 0; r < num_rows; ++r) {
        auto mode = data.row(r)
                    .sort()
                    .rle()
                    .max_by_key(parrot::snd())
                    .value();
        auto index = data.row(r).last_index_of(mode);
        results.push_back(thrust::make_pair(mode, index));
    }

    return parrot::array(results);
}
```



Parrot

A high level, parallel, array-based library
with implicit fusion

Thank You

<https://github.com/codereport/Content/Talks>

Conor Hoekstra



code_report



codereport

Questions?

<https://github.com/codereport/Content/Talks>
choekstra@nvidia.com

Conor Hoekstra



code_report



codereport