

Functional GPU Programming

Conor Hoekstra



code_report | codereport



Array GPU Programming

Conor Hoekstra



code_report | codereport





Programming Language Rankings (2025 Aug)

by code_report





StackOverflow Octoverse RedMonk Languish JetBrains
 IEEE Spectrum PYPL TIOBE GitHub 2.0 [Rankings Overview](#)

Exclude "Edge Languages" | Number of Languages: [20](#) Months for Delta (Δ): [3](#) | [All](#) Languages

		Language	Avg	StDev	n ¹	3mΔ		Language	Avg	StDev	n ¹	3mΔ	
1		JavaScript	1.5	0.57	4	-	11		Rust	12.66	3.21	3	-
2		Python	2	1.41	4	-	12		Kotlin	14	1.73	3	-
3		TypeScript	4.5	1.73	4	-	13		Swift	15.66	5.13	3	-
4		Java	4.5	1.73	4	-	14		PowerShell	15.66	6.65	3	(1)
5		C#	5.75	1.5	4	-	15		R	16	5	3	(1)
6		C++	7	1.41	4	-	16		Dart	17.66	2.3	3	-
7		PHP	8.5	3.69	4	-	17		Ruby	18	10	3	-
8		Shell	9.5	3.69	4	(1)	18		Lua	21.5	7.77	2	-
9		C	9.75	0.95	4	(1)	19		Objective-C	23	14.14	2	(2)
10		Go	10.25	2.5	4	-	20		VBA	25	1.41	2	(1)

1 - The number of (selected) ranking websites this language shows up in.

If you have suggestions or find a bug, you can open an [issue](#) here.






383 Videos

48 (32) Talks

Algorithms +
Data
Structures =
Programs



TACIT
Ξ—ΛVV—Ξ
TALK

254 Episodes
@adspthepodcast



115 Episodes
@arraycast



27 Episodes
@codereport



Home

Shorts

Subscriptions

YouTube Mu...

You

Downloads




code_report

@code_report · 60.8K subscribers · 352 videos

Welcome to the code_report YouTube channel. ...[more](#)twitter.com/code_report and 3 more links[Customize channel](#)[Manage videos](#)
[Home](#) [Videos](#) [Live](#) [Podcasts](#) [Playlists](#) [Community](#)
[Latest](#)[Popular](#)[Oldest](#)

1 Problem, 24 Programming Languages

375K views · 1 year ago



From C → C++ → Rust

170K views · 1 year ago

1 Problem, 16 Programming Languages
(C++ vs Rust vs Haskell vs Python vs APL...)

158K views · 3 years ago



Functional vs Array Programming

131K views · 3 years ago



<https://github.com/codereport/Content>





Parrot



Parrot

A high level, parallel, array-based library
with implicit fusion



Parrot

<https://github.com/nvlabs/parrot>



Parrot

<https://github.com/nvlabs/parrot>



Parrot

<https://github.com/nvlabs/parrot>

Apache 2.0 License

Overview

Problem 1



API Overview



Problem 2



Problem 3



Sum of Squares (SOS)






```
#include <iostream>
#include <numeric>
#include <ranges>

auto sos(int N) {
    return std::ranges::views::iota(0, N) |
        std::ranges::views::transform([](int x) { return x * x; }) |
        std::ranges::fold_left(0, std::plus{});
}

int main() { std::cout << sos(10) << std::endl; }
```



```
#include <algorithm>
#include <iostream>
#include <numeric>
#include <ranges>

auto sos(int N) {
    return std::ranges::views::iota(0, N) |
        std::ranges::views::transform([](int x) { return x * x; }) |
        std::ranges::fold_left(0, std::plus{});
}

int main() { std::cout << sos(10) << std::endl; }
```



```
#include <algorithm>
#include <iostream>
#include <ranges>

auto sos(int N) {
    return std::ranges::views::iota(0, N) |
        std::ranges::views::transform([](int x) { return x * x; }) |
        std::ranges::fold_left(0, std::plus{});
}

int main() { std::cout << sos(10) << std::endl; }
```



```
#include <algorithm>
#include <iostream>
#include <ranges>

auto sos(int N) {
    return std::ranges::fold_left(
        std::ranges::views::iota(0, N) |
        std::ranges::views::transform([](int x) { return x * x; }),
        0,
        std::plus{});
}

int main() { std::cout << sos(10) << std::endl; }
```



```
#include <algorithm>
#include <print>
#include <ranges>

auto sos(int N) {
    return std::ranges::fold_left(
        std::ranges::views::iota(0, N) |
        std::ranges::views::transform([](int x) { return x * x; }),
        0,
        std::plus{});
}

int main() { std::print("{}", sos(10)); }
```





```
#include <thrust/device_vector.h>
#include <thrust/functional.h>
#include <thrust/host_vector.h>
#include <thrust/reduce.h>
#include <thrust/sequence.h>
#include <thrust/transform.h>
#include <iostream>

// Functor to square a number
struct square {
    __host__ __device__ int operator()(const int& x) const { return x * x; }
};

auto sos(int N) {
    // Create a device vector and fill it with sequence 0, 1, 2, ..., N-1
    thrust::device_vector<int> d_vec(N);
    thrust::sequence(d_vec.begin(), d_vec.end());

    // Square each element and sum the result
    return thrust::transform_reduce(
        d_vec.begin(), d_vec.end(), square(), 0, thrust::plus<int>());
}

int main() {
    auto result = sos(10);
    std::cout << result << std::endl;
    return 0;
}
```



```
#include <thrust/device_vector.h>
#include <thrust/functional.h>
#include <thrust/host_vector.h>
#include <thrust/reduce.h>
#include <thrust/sequence.h>
#include <thrust/transform_reduce.h>
#include <iostream>

// Functor to square a number
struct square {
    __host__ __device__ int operator()(const int& x) const { return x * x; }
};

auto sos(int N) {
    // Create a device vector and fill it with sequence 0, 1, 2, ..., N-1
    thrust::device_vector<int> d_vec(N);
    thrust::sequence(d_vec.begin(), d_vec.end());

    // Square each element and sum the result
    return thrust::transform_reduce(
        d_vec.begin(), d_vec.end(), square(), 0, thrust::plus<int>());
}

int main() {
    auto result = sos(10);
    std::cout << result << std::endl;
    return 0;
}
```



```
#include <thrust/device_vector.h>
#include <thrust/sequence.h>
#include <thrust/transform_reduce.h>
#include <iostream>

// Functor to square a number
struct square {
    __host__ __device__ int operator()(const int& x) const { return x * x; }
};

auto sos(int N) {
    // Create a device vector and fill it with sequence 0, 1, 2, ..., N-1
    thrust::device_vector<int> d_vec(N);
    thrust::sequence(d_vec.begin(), d_vec.end());

    // Square each element and sum the result
    return thrust::transform_reduce(
        d_vec.begin(), d_vec.end(), square(), 0, thrust::plus<int>());
}

int main() {
    auto result = sos(10);
    std::cout << result << std::endl;
    return 0;
}
```



```
#include <thrust/device_vector.h>
#include <thrust/sequence.h>
#include <thrust/transform_reduce.h>
#include <iostream>

struct square {
    __host__ __device__ int operator()(const int& x) const { return x * x; }
};

auto sos(int N) {
    thrust::device_vector<int> d_vec(N);
    thrust::sequence(d_vec.begin(), d_vec.end());
    return thrust::transform_reduce(
        d_vec.begin(), d_vec.end(), square(), 0, thrust::plus<int>());
}

int main() {
    auto result = sos(10);
    std::cout << result << std::endl;
    return 0;
}
```



```
#include <thrust/device_vector.h>
#include <thrust/sequence.h>
#include <thrust/transform_reduce.h>
#include <iostream>

auto sos(int N) -> int {
    thrust::device_vector<int> d_vec(N);
    thrust::sequence(d_vec.begin(), d_vec.end());
    return thrust::transform_reduce(
        d_vec.begin(),
        d_vec.end(),
        [] __host__ __device__(int x) { return x * x; },
        0,
        thrust::plus<int>());
}

int main() {
    auto result = sos(10);
    std::cout << result << std::endl;
    return 0;
}
```



```
#include <thrust/device_vector.h>
#include <thrust/sequence.h>
#include <thrust/transform_reduce.h>
#include <iostream>

auto sos(int N) -> int {
    auto iota = thrust::counting_iterator<int>(0);
    return thrust::transform_reduce(
        iota,
        iota + N,
        [] __host__ __device__(int x) { return x * x; },
        0,
        thrust::plus<int>());
}

int main() {
    auto result = sos(10);
    std::cout << result << std::endl;
    return 0;
}
```



```
#include <thrust/iterator/counting_iterator.h>
#include <thrust/iterator/transform_iterator.h>
#include <thrust/reduce.h>
#include <iostream>

auto sos(int N) -> int {
    auto iota = thrust::counting_iterator<int>(0);
    auto map = thrust::make_transform_iterator(
        iota, [] __host__ __device__(int x) { return x * x; });
    return thrust::reduce(map, map + N, 0);
}

int main() {
    auto result = sos(10);
    std::cout << result << std::endl;
    return 0;
}
```





```
auto sos(int N) {  
}
```



```
auto sos(int N) {  
    return parrot::range(N)  
}
```



```
auto sos(int N) {  
    return parrot::range(N)  
        .sq()  
}
```



```
auto sos(int N) {  
    return parrot::range(N)  
        .sq()  
        .sum();  
}
```



```
auto sos(int N) {  
    return parrot::range(N).print()  
        .sq().print()  
        .sum().print();  
}
```



```
#include "parrot.hpp"

auto sos(int N) {
    return parrot::range(N)
        .sq()
        .sum();
}

int main() { sos(10).print(); }
```



```
#include "parrot.hpp"

auto sos(int N) {
    return parrot::range(N).sq().sum();
}

int main() { sos(10).print(); }
```



Parrot API

Fused

1-index Maps (Unary)

[map](#)
[abs](#)
[dble](#)
[enumerate](#)
[even](#)
[half](#)
[log](#)
[exp](#)
[neg](#)
[odd](#)
[rand](#)
[sign](#)
[sq](#)
[sqrt](#)

1-index Maps (Binary)

[map2](#)
[add \(+\)](#)
[div \(/\)](#)
[gt \(>\)](#)
[gte \(>=\)](#)
[idiv](#)
[lt \(<\)](#)
[lte \(<=\)](#)
[max](#)
[min](#)
[minus \(-\)](#)
[times \(*\)](#)
[eq \(==\)](#)
[pairs](#)

2-index Maps

[map_adj](#)
[deltas](#)
[differ](#)
Joins
[append](#)
[prepend](#)
Products
[cross](#)
[outer](#)
Reshapes
[take](#)
[drop](#)
[transpose](#)
[reshape](#)
[cycle](#)
[repeat](#)
Copying
[replicate](#)
Permutations
[rev](#)
[gather](#)

Conditionally Fused

Compactions
[keep](#)
[filter](#)
[where](#)
[uniq](#)
[distinct](#)

Materializing

Reductions

[reduce](#)
[all](#)
[any](#)
[maxr](#)
[max_by_key](#)
[minr](#)
[minmax](#)
[prod](#)
[sum](#)
Scans
[scan](#)
[alls](#)
[anyss](#)
[maxs](#)
[mins](#)
[prods](#)
[sums](#)

Permutations

[sort](#)
[sort_by](#)
[sort_by_key](#)

Compactions

[rle](#)
Copying
[replicate](#)
Split-Reductions
[chunk_by_reduce](#)
Comparisons
[match](#)

Properties

[size](#)
[rank](#)
[shape](#)

Accessors

[value](#)
[front](#)
[back](#)
[to_host](#)

Array Creation

[array](#)
[range](#)
[scalar](#)
[matrix](#)

[I/O](#)
[print](#)

Function Objects

Accessors
[fst](#)
[snd](#)

Binary Operations


[eq](#)
[gt](#)
[gte](#)
[lt](#)
[lte](#)
[max](#)
[min](#)
[mul](#)
[add](#)

Functional vs Array Programming

Array = Functional
+ Rank Polymorphism

Array* = Functional
+ Rank Polymorphism
+ Symbols


* Iversonian (APL, BQN, Iua)



Functional

Array

Iversonian
Array



Fused

1-index Maps (Unary)

[map](#)
[abs](#)
[dble](#)
[enumerate](#)
[even](#)
[half](#)
[log](#)
[exp](#)
[neg](#)
[odd](#)
[rand](#)
[sign](#)
[sq](#)
[sqrt](#)

1-index Maps (Binary)

[map2](#)
[add \(+\)](#)
[div \(/\)](#)
[gt \(>\)](#)
[gte \(>=\)](#)
[idiv](#)
[lt \(<\)](#)
[lte \(<=\)](#)
[max](#)
[min](#)
[minus \(-\)](#)
[times \(*\)](#)
[eq \(==\)](#)
[pairs](#)

2-index Maps

[map_adj](#)
[deltas](#)
[differ](#)
Joins
[append](#)
[prepend](#)
Products
[cross](#)
[outer](#)
Reshapes
[take](#)
[drop](#)
[transpose](#)
[reshape](#)
[cycle](#)
[repeat](#)

Copying
[replicate](#)
Permutations
[rev](#)
[gather](#)

Conditionally Fused

Compactions
[keep](#)
[filter](#)
[where](#)
[uniq](#)
[distinct](#)

Materializing

Reductions

[reduce](#)
[all](#)
[any](#)
[maxr](#)
[max_by_key](#)
[minr](#)
[minmax](#)
[prod](#)
[sum](#)
Scans
[scan](#)
[alls](#)
[anyss](#)
[maxs](#)
[mins](#)
[prods](#)
[sums](#)

Permutations

[sort](#)
[sort_by](#)
[sort_by_key](#)

Compactions

[rle](#)
Copying
[replicate](#)

Split-Reductions

[chunk_by_reduce](#)
Comparisons
[match](#)

Properties

[size](#)
[rank](#)
[shape](#)

Accessors

[value](#)
[front](#)
[back](#)
[to_host](#)

Array Creation

[array](#)
[range](#)
[scalar](#)
[matrix](#)

I/O

[print](#)

Function Objects

Accessors
[fst](#)
[snd](#)

Binary Operations

[eq](#)
[gt](#)
[gte](#)
[lt](#)
[lte](#)
[max](#)
[min](#)
[mul](#)
[add](#)

Fused
1-index Maps (Unary)

 [map](#)

[abs](#)

[dble](#)

[enumerate](#)

[even](#)

[half](#)

[log](#)

[exp](#)

[neg](#)

[odd](#)

[rand](#)

[sign](#)

[sq](#)

[sqrt](#)

Fused
1-index Maps (Binary)

 [map2](#)

[add \(+\)](#)

[div \(/\)](#)

[gt \(>\)](#)

[gte \(>=\)](#)

[idiv](#)

[lt \(<\)](#)

[lte \(<=\)](#)

[max](#)


[min](#)

[minus \(-\)](#)

[times \(*\)](#)

[eq \(==\)](#)

[pairs](#)



```
auto add_one(auto& matrix) {
    for (auto& row : matrix) {
        for (auto& e : row) {
            e += 1;
        }
    }
    return matrix;
}
```



```
auto add_one(const auto& matrix) {
    return matrix | std::ranges::views::transform([](const auto& row) {
        return row | std::ranges::views::transform(
            [](auto e) { return e + 1; });
    });
}
```



```
auto add_one(auto matrix) {  
    return matrix + 1;  
}
```



```
auto add_one(auto matrix) {  
    return matrix.add(1);  
}
```



```
auto add_one(auto matrix) {  
    return matrix + 1;  
}
```

Fused
1-index Maps (Unary)

 [map](#)

[abs](#)

[dble](#)

[enumerate](#)

[even](#)

[half](#)

[log](#)

[exp](#)

[neg](#)

[odd](#)

[rand](#)

[sign](#)

[sq](#)

[sqrt](#)

Fused
1-index Maps (Binary)

 [map2](#)

[add \(+\)](#)

[div \(/\)](#)

[gt \(>\)](#)

[gte \(>=\)](#)

[idiv](#)

[lt \(<\)](#)

[lte \(<=\)](#)

[max](#)

[min](#)

[minus \(-\)](#)

[times \(*\)](#)

[eq \(==\)](#)

[pairs](#)

Fused

1-index Maps (Unary)

[map](#)
[abs](#)
[dble](#)
[enumerate](#)
[even](#)
[half](#)
[log](#)
[exp](#)
[neg](#)
[odd](#)
[rand](#)
[sign](#)
[sq](#)
[sqrt](#)

1-index Maps (Binary)

[map2](#)
[add \(+\)](#)
[div \(/\)](#)
[gt \(>\)](#)
[gte \(>=\)](#)
[idiv](#)
[lt \(<\)](#)
[lte \(<=\)](#)
[max](#)
[min](#)
[minus \(-\)](#)
[times \(*\)](#)
[eq \(==\)](#)
[pairs](#)

2-index Maps

[map_adj](#)
[deltas](#)
[differ](#)
Joins
[append](#)
[prepend](#)
Products
[cross](#)
[outer](#)
Reshapes
[take](#)
[drop](#)
[transpose](#)
[reshape](#)
[cycle](#)
[repeat](#)
Copying
[replicate](#)
Permutations
[rev](#)
[gather](#)

Conditionally Fused

Compactions
[keep](#)
[filter](#)
[where](#)
[uniq](#)
[distinct](#)

Materializing

Reductions

[reduce](#)
[all](#)
[any](#)
[maxr](#)
[max_by_key](#)
[minr](#)
[minmax](#)
[prod](#)
[sum](#)
Scans
[scan](#)
[alls](#)
[anyss](#)
[maxs](#)
[mins](#)
[prods](#)
[sums](#)

Permutations

[sort](#)
[sort_by](#)
[sort_by_key](#)

Compactions

[rle](#)
Copying
[replicate](#)
Split-Reductions
[chunk_by_reduce](#)
Comparisons
[match](#)

Properties

[size](#)
[rank](#)
[shape](#)

Accessors

[value](#)
[front](#)
[back](#)
[to_host](#)

Array Creation

[array](#)
[range](#)
[scalar](#)
[matrix](#)

I/O

[print](#)

Function Objects

Accessors
[fst](#)
[snd](#)

Binary Operations

[eq](#)
[gt](#)
[gte](#)
[lt](#)
[lte](#)
[max](#)
[min](#)
[mul](#)
[add](#)

Problem 2

(taken from PaddlePaddle)



```
// https://github.com/PaddlePaddle/Paddle/blob/80f1123eb0c...
```

```
template <typename T>
static void Get_____(  
    const phi::GPUContext& dev_ctx,  
    const DenseTensor* input_tensor,  
    const int64_t num_cols,  
    const int64_t num_rows,  
    T* out_tensor,  
    int64_t* indices_tensor)  
{  
    DenseTensor input_tmp;  
    input_tmp.Resize(common::make_ddim({num_rows, num_cols}));  
    T* input_tmp_data = dev_ctx.Alloc<T>(&input_tmp);  
    phi::Copy(dev_ctx, *input_tensor, dev_ctx.GetPlace(), false, &input_tmp);  
  
    thrust::device_ptr<T> out_tensor_ptr(out_tensor);  
    thrust::device_ptr<int64_t> indices_tensor_ptr(indices_tensor);  
  
    for (int64_t i = 0; i < num_rows; ++i) {  
        T* begin = input_tmp_data + num_cols * i;  
        T* end = input_tmp_data + num_cols * (i + 1);  
        thrust::device_vector<int64_t> indices_data(num_cols);  
        thrust::sequence(thrust::device,  
            indices_data.begin(),  
            indices_data.begin() + num_cols);  
        thrust::sort_by_key(thrust::device, begin, end, indices_data.begin());  
        int unique = 1 + thrust::inner_product(thrust::device,  
            begin,  
            begin,  
            end - 1,  
            begin + 1,  
            0,  
            thrust::plus<int>(),  
            thrust::not_equal_to<T>());  
        thrust::device_vector<T> keys_data(unique);  
        thrust::device_vector<int64_t> cnts_data(unique);  
        thrust::reduce_by_key(thrust::device,  
            begin,  
            end,  
            thrust::constant_iterator<int>(1),  
            keys_data.begin(),  
            cnts_data.begin());  
        auto it = thrust::max_element(  
            thrust::device, cnts_data.begin(), cnts_data.begin() + unique);  
        T ____ = keys_data[it - cnts_data.begin()];  
        int64_t counts = cnts_data[it - cnts_data.begin()];  
        auto pos = thrust::find(thrust::device, begin, end, mode);  
        int64_t index = indices_data[pos - begin + counts - 1];  
        out_tensor_ptr[i] = static_cast<T>(mode);  
        indices_tensor_ptr[i] = static_cast<int64_t>(index);  
    }  
}
```



```
// https://github.com/PaddlePaddle/Paddle/blob/80f1123eb0c...

template <typename T>
static void Get_____(/* ... */) {
    // initialization

    for (int64_t i = 0; i < num_rows; ++i) {
        T* begin = input_tmp_data + num_cols * i;
        T* end = input_tmp_data + num_cols * (i + 1);
        thrust::device_vector<int64_t> indices_data(num_cols);
        thrust::sequence(thrust::device,
            indices_data.begin(),
            indices_data.begin() + num_cols);
        thrust::sort_by_key(thrust::device, begin, end, indices_data.begin());
        int unique = 1 + thrust::inner_product(thrust::device,
            begin,
            end - 1,
            begin + 1,
            0,
            thrust::plus<int>(),
            thrust::not_equal_to<T>());
        thrust::device_vector<T> keys_data(unique);
        thrust::device_vector<int64_t> cnts_data(unique);
        thrust::reduce_by_key(thrust::device,
            begin,
            end,
            thrust::constant_iterator<int>(1),
            keys_data.begin(),
            cnts_data.begin());
        auto it = thrust::max_element(
            thrust::device, cnts_data.begin(), cnts_data.begin() + unique);
        T ____ = keys_data[it - cnts_data.begin()];
        int64_t counts = cnts_data[it - cnts_data.begin()];
        auto pos = thrust::find(thrust::device, begin, end, mode);
        int64_t index = indices_data[pos - begin + counts - 1];
        out_tensor_ptr[i] = static_cast<T>(mode);
        indices_tensor_ptr[i] = static_cast<int64_t>(index);
    }
}
```



```
int unique = 1 + thrust::inner_product(
    begin,
    end - 1,
    begin + 1,
    0,
    thrust::plus<int>(),
    thrust::not_equal_to<T>());
```



```
int unique = 1 + thrust::inner_product(
    begin,
    end - 1,
    begin + 1,
    0,
    thrust::plus<int>(),
    thrust::not_equal_to<T>());
```



```
int unique = thrust::unique_count(begin, end);
```

[[digression]]





```
template <typename I>
auto unique_count(I first, I last) {
    auto zip = thrust::make_zip_iterator(first, first + 1);
    auto map = thrust::make_transform_iterator(
        zip, thrust::make_zip_function(thrust::not_equal_to{}));
    return 1 + thrust::reduce(map,
                            map + (last - first - 1),
                            0,
                            thrust::plus{});
}
```



```
template <class Derived,
          class ForwardIt,
          class BinaryPred>
typename thrust::iterator_traits<ForwardIt>::difference_type
_CCCL_HOST_DEVICE
unique_count(execution_policy<Derived> &policy,
             ForwardIt first,
             ForwardIt last,
             BinaryPred binary_pred)
{
    if (first == last) {
        return 0;
    }
    auto size = thrust::distance(first, last);
    auto it   = thrust::make_zip_iterator(thrust::make_tuple(first, thrust::next(first)));
    return 1 + thrust::count_if(policy, it, thrust::next(it, size - 1),
                               zip_adj_not_predicate<BinaryPred>{binary_pred});
}
```



```
template <typename I>
auto unique_count_zip_count_if(I first, I last) {
    auto zip = thrust::make_zip_iterator(first, first + 1);
    auto neq = thrust::make_zip_function(thrust::not_equal_to{});
    return 1 + thrust::count_if(zip, zip + (last - first - 1), neq);
}
```



```
auto unique_count(auto data) {
    return data.map_adj(parrot::neq{}).sum() + 1;
}
```



```
auto unique_count(auto data) {  
    return data.differ().sum() + 1;  
}
```



thrust::transform_reduce

transform_iterator + reduce

★ map + reduce

thrust::unique_count

zip_iterator + count_if

★ map_adj + sum

thrust::tabulate

counting_iterator +
transform

range + ★ map

[[end of digression]]



```
// https://github.com/PaddlePaddle/Paddle/blob/80f1123eb0c...

template <typename T>
static void Get_____/* ... */ {
    // initialization

    for (int64_t i = 0; i < num_rows; ++i) {
        T* begin = input_tmp_data + num_cols * i;
        T* end = input_tmp_data + num_cols * (i + 1);
        thrust::device_vector<int64_t> indices_data(num_cols);
        thrust::sequence(thrust::device,
                        indices_data.begin(),
                        indices_data.begin() + num_cols);
        thrust::sort_by_key(thrust::device, begin, end, indices_data.begin());
        int unique = thrust::unique_count(thrust::device, begin, end);
        thrust::device_vector<T> keys_data(unique);
        thrust::device_vector<int64_t> cnts_data(unique);
        thrust::reduce_by_key(thrust::device,
                            begin,
                            end,
                            thrust::constant_iterator<int>(1),
                            keys_data.begin(),
                            cnts_data.begin());
        auto it = thrust::max_element(
            thrust::device, cnts_data.begin(), cnts_data.begin() + unique);
        T _____ = keys_data[it - cnts_data.begin()];
        int64_t counts = cnts_data[it - cnts_data.begin()];
        auto pos = thrust::find(thrust::device, begin, end, mode);
        int64_t index = indices_data[pos - begin + counts - 1];
        out_tensor_ptr[i] = static_cast<T>(mode);
        indices_tensor_ptr[i] = static_cast<int64_t>(index);
    }
}
```



```
// https://github.com/PaddlePaddle/Paddle/blob/80f1123eb0c...

template <typename T>
static void Get_____/* ... */ {
    // initialization

    for (int64_t i = 0; i < num_rows; ++i) {
        T* begin      = input_tmp_data + num_cols * i;
        T* end       = input_tmp_data + num_cols * (i + 1);
        auto indices_data = thrust::device_vector<int64_t>(num_cols);

        thrust::sequence(indices_data.begin(), indices_data.begin() + num_cols);
        thrust::sort_by_key(begin, end, indices_data.begin());

        int unique     = thrust::unique_count(thrust::device, begin, end);
        auto keys_data = thrust::device_vector<T>(unique);
        auto cnts_data = thrust::device_vector<int64_t>(unique);

        thrust::reduce_by_key(
            begin,
            end,
            thrust::constant_iterator<int>(1),
            keys_data.begin(),
            cnts_data.begin());

        auto it = thrust::max_element(cnts_data.begin(), cnts_data.begin() + unique);

        T _____ = keys_data[it - cnts_data.begin()];
        int64_t counts      = cnts_data[it - cnts_data.begin()];
        auto pos          = thrust::find(begin, end, mode);
        int64_t index      = indices_data[pos - begin + counts - 1];
        out_tensor_ptr[i]   = static_cast<T>(mode);
        indices_tensor_ptr[i] = static_cast<int64_t>(index);
    }
}
```



```
template <typename Array>
auto GetModeBySort_Parrot(
    const Array& data, int num_rows, int num_cols) {
using T = typename Array::value_type;
std::vector<thrust::pair<T, int>> results;

for (int r = 0; r < num_rows; ++r) {
    auto mode = parrot::stats::mode(data.row(r)).value();
    auto index = data.row(r).last_index_of(mode);
    results.push_back(thrust::make_pair(mode, index));
}

return parrot::array(results);
}
```



```
template <typename Array>
auto GetModeBySort_Parrot(
    const Array& data, int num_rows, int num_cols) {
    using T = typename Array::value_type;
    std::vector<thrust::pair<T, int>> results;

    for (int r = 0; r < num_rows; ++r) {
        auto mode = data.row(r)
                    .sort()
                    .rle()
                    .max_by_key(parrot::snd())
                    .value();
        auto index = data.row(r).last_index_of(mode);
        results.push_back(thrust::make_pair(mode, index));
    }

    return parrot::array(results);
}
```

Softmax




```
int main() {
    auto matrix = // ... (matrix initialization)

    auto cols = matrix.shape()[1];
    auto z    = matrix - matrix.maxr<2>().replicate(cols);
    auto num  = z.exp();
    auto den  = num.sum<2>();
    (num / den.replicate(cols)).print();

    return 0;
}
```

kp

CUDA Kernel Profiling (softmax)





Parrot


A high level, parallel, array-based library
with implicit fusion



Parrot

<https://github.com/nvlabs/parrot>

Apache 2.0 License



Thank You

<https://github.com/codereport/Content/Talks>

Conor Hoekstra

 code_report

 codereport

Free Instructor-Led CUDA C++ Training

As a C++ Under the Sea attendee, you have exclusive, free access to a full-day, virtual XLab. Join NVIDIA experts on October 22 as they explain how to build a CUDA application in C++, from start to finish.

Spaces are limited so register today to secure your place.



[Register Now](#)





Questions?



<https://github.com/codereport/Content/Talks>
choekstra@nvidia.com

Conor Hoekstra


 code_report


 codereport

Sushi for Two



<https://codeforces.com/contest/1138/problem/A>











```
template <int N>
constexpr auto sushi_for_two(std::array<int, N> sushi) {
    int current_sushi      = 0;
    int sushi_in_a_row     = 0;
    int prev_sushi_in_a_row = 0;
    int max_of_mins        = 0;
    for (auto const s : sushi) {
        if (current_sushi != s) {
            current_sushi = s;
            if (prev_sushi_in_a_row == 0) {
                prev_sushi_in_a_row = sushi_in_a_row;
                sushi_in_a_row     = 1;
            } else {
                auto const min      = std::min(sushi_in_a_row, prev_sushi_in_a_row);
                max_of_mins        = std::max(max_of_mins, min);
                prev_sushi_in_a_row = sushi_in_a_row;
                sushi_in_a_row     = 1;
            }
        } else {
            sushi_in_a_row += 1;
        }
    }
    auto const min = std::min(sushi_in_a_row, prev_sushi_in_a_row);
    max_of_mins   = std::max(max_of_mins, min);
    return max_of_mins * 2;
}
```



```
using namespace ranges::views;
using namespace combinator;

auto sushi_for_two(std::vector<int> sushi) {
    auto indices = concat(
        concat(single(0),
            zip(zip_with(_neq_, sushi, sushi | drop(1)),
                iota(1))
            | filter(_fst)
            | values),
        single(sushi.size())));
    auto deltas = zip_with(_sub_, indices | drop(1), indices);
    return 2 * ranges::max(zip_with(_min_, deltas, deltas | drop(1)));
}
```



```
auto sushi_for_two(auto sushi) {
    return sushi.differ()
        .where()
        .prepend(0)
        .append(sushi.size())
        .deltas()
        .map_adj(parrot::min{})
        .dble()
        .maxr();
}
```